

Function pointers, Signals and low level file I/O

COMP2017/COMP9017





- › So far all variables are exposed as having an address
- › Compiled binary code is no different

```
if ( != 0 ) {  
    execute statement;  
} else {  
    execute other statement;  
}
```



Function Pointers

```
int x = 33;  
if (x == 33)  
{  
    x = 480+7;  
}  
x += 768;
```



```
movl $33, -4(%rbp)  
cmpl $33, -4(%rbp)  
jne .L2      jump not equal to  
movl $487, -4(%rbp)  
.L2:  
    addl $768, -4(%rbp)
```

- › JUMP!
- › Same with loops
- › rbp is the stack frame pointer on x86_64



Function Pointers

```
int x = 33;  
if (x == 33)  
{  
    x = 480+7;  
    foo();  
}  
x += 768;
```



```
movl $33, -4(%rbp)  
cmpl $33, -4(%rbp)  
jne .L4  
movl $487, -4(%rbp)  
movl $0, %eax  
call foo  
.L4:  
addl $768, -4(%rbp)
```

- › Call? If not a jump, how do we get back?



Function Pointers

```
int x = 33;  
if (x == 33)  
{  
    x = 480+7;  
    foo();  
}  
x += 768;
```



```
movl $33, -4(%rbp)  
cmpl $33, -4(%rbp)  
jne .L4  
movl $487, -4(%rbp)  
movl $0, %eax  
call foo  
.L4:  
addl $768, -4(%rbp)
```

- › Call? If not a jump, how do we get back?
- › Stack is being managed here. Callee or caller will setup and teardown the stack



Function Pointers

```
int (*fptr) () = foo;
```

assign the function pointer
address to the variable fptr

```
int x = 33;  
if (x == 33) {  
    x = 480+7;  
    fptr();  
}  
x += 768;
```

we change the
implementation of function
during runtime



```
subq $16, %rsp  
leaq foo(%rip), %rax  
movq %rax, -16(%rbp)  
movl $33, -4(%rbp)  
cmpl $33, -4(%rbp)  
jne .L4  
...
```

- › If we jump, or call, all we need is an address

a reference to some area of memory which we may have a function to call to, the pointer can take value dynamically

```
origin: void memcpy(void* dst, void* src, size_t size);  
declaration: void (*my_memcpy)(void*, void*, size_t) = memcpy;  
call: my_memcpy(buffer, src_buffer, 100);
```



Function Pointers

```
int (*fptr) () = foo;  
  
int x = 33;  
if (x == 33)  
{  
    x = 480+7;  
    fptr();  
}  
x += 768;
```



```
...  
movl $33, -4(%rbp)  
cmpl $33, -4(%rbp)  
jne .L4  
movl $487, -4(%rbp)  
movq -16(%rbp), %rdx  
movl $0, %eax  
call *%rdx  
.L4:  
addl $768, -4(%rbp)
```

- › Call a function, jump to address is (almost) the same process



- › A function pointer is an address that refers to an area of memory with executable code
- › Typically the first instruction of the function call[^]
- › Are useful for conventional programming patterns
- › Examples
 - Do something, and when you are finished call this function
 - Do something, and if it goes wrong, call this function
 - I am a data source, give me an function to send new bits of data to
 - I want to sort a list of objects, here is the address of a function to perform comparison of two elements

[^] Depends on callee/caller conventions



- › The declaration of the function pointer parameter looks like:
`type (*f)(param declaration...)`

- › and the call of the function looks like:
`f(params...)`



Function pointer: Example

- › Call functionA if x is true, or functionB otherwise

```
if (x)
    funcA();
else
    funcB();
```



Function pointer: Example

- › Call functionA if x is true, or functionB otherwise

```
if (x)
    funcA();
else
    funcB();
```

- › What if we don't know what funcA and funcB are at compile time?

```
void do_process(int x, funcA?, funcB?) {
    if (x)
        funcA(x); // print x
    else
        funcB(x); // delete elem x
}
```



Function pointer: Example

- › What if we don't know what funcA and funcB are at compile time?

```
void deleteX(int x);
void printX(int x);

void do_process(int x,
    void __funcA__(int),
    void __funcB__(int))
{
    if (x)
        funcA(x); // print X
    else
        funcB(x); // delete elem X
}
```



Function pointer: Example

- › What if we don't know what funcA and funcB are at compile time?

```
void deleteX(int x);
void printX(int x);

void do_process(int x,
    void (*funcA)(int),
    void (*funcB)(int))
{
    if (x)
        funcA(x); // print X
    else
        funcB(x); // delete elem X
}
```



Function pointer: Example

- › Write less code. Allow option to change implementation choices at runtime. E.g. heuristics, look and feel, plugins

```
void printX_1(int x) { printf("%d", x); }
void printX_2(int x) { printf("%d\n", x); }
void printX_3(int x) { printf("x: %d\n", x); }
```

```
// delegate which fn pointer
if (user_style == PRETTY)
    print_style = printX_3;
...
// generic code
do_process(value1, print_style, remove_style);
do_process(value2, print_style, remove_style);
do_process(value3, print_style, remove_style);
```

Signals

COMP2017/COMP9017

FACULTY OF
ENGINEERING



THE UNIVERSITY OF
SYDNEY





- › a process can communicate with another using a *signal*
- › these are a form of *software interrupt*
- › execution is interrupted and a function call is made at that point to a user specified function
- › when the function returns, execution is resumed restores

when we stop the program, we need to save the state on what is happening

```
#include <stdio.h>
#include <signal.h>
#include <unistd.h>

volatile int interrupted = 0;

void impatient(int arg) {           Signal Handler
    interrupted = 1;
}

int main() {
    void (*variable)(int) = impatient;  function pointer
    signal(SIGINT, variable); if the SIGINT, call variable. SIGINT stands for signal interrupted
    printf("We are going to do something for a long time\n");

    while (!interrupted)      ctl + c to send SIGINT to ./a.out
        usleep(10);

    printf("Oh, you didn't like waiting\n");
    printf("Program terminated\n");

    return 0;
}
```



- › signals can be generated by one process to another using the *kill* system call
 - to get <pid>: \$ ps -A | grep a.out
- › signals are also generated by the operating system, eg when an access outside memory bounds is attempted (Segmentation Fault)

1. \$ kill -9 <pid>
2. \$ kill -s SIGKILL <pid>
3. \$ kill -s SIGSEGV <pid>



SIGHUP	1 Hangup	SIGKILL	9 Kill
SIGINT	2 Interrupt	SIGBUS	10 Bus Error
SIGQUIT	3 Quit	SIGSEGV	11 Segmentation Fault
SIGILL	4 Illegal Instruction	SIGSYS	12 Bad System Call
SIGTRAP	5 Trace or Breakpoint Trap	SIGPIPE	13 Broken Pipe
SIGABRT	6 Abort	SIGALRM	14 Alarm Clock
SIGEMT	7 Emulation Trap	SIGTERM	15 Terminated
SIGFPE	8 Arithmetic Exception	SIGUSR1	16 User Signal 1
		SIGUSR2	17 User Signal 2



- › You can send a signal to a running process from the command line using the kill command

› Eg `kill -9 12345`

Will send the SIGKILL signal to process 12345.

- › Some signals can be *caught* and handled by a user supplied function
- › Some signals (such as SIGKILL) cannot be caught and caused the process to be terminated

```
1 #include <stdio.h>
2 #include <signal.h>
3 #include <unistd.h>
4 #include <signal.h>
5 #include <errno.h>
6
7 volatile int flag = 0;
8
9 // signal handler
10 void interrupted(int arg) {
11     flag = 1;
12 }
13
14 int main() {
15     void (*old_sig_int_handler)(int);
16     int res;
17
18     // get the old handler
19     old_sig_int_handler = signal(SIGINT, interrupted);
20     // overwrite SIGINT
21     if (old_sig_int_handler == SIG_ERR) {
22         perror("could not change signal handler");
23         return -1;
24     }
25
26     char buffer[100];
27     flag = 0;
28     // read from standard input
29     ssize_t result = read(0, buffer, 100);
30 }
```

```
31     // check for errors
32     int error_val = errno;
33     if (error_val != 0) {
34         printf("\n");
35         printf("error_val: %d\n", error_val);
36         printf("read() was interrupted by a signal\n");
37         printf("flag is: %d\n", flag);
38         perror("hmm errno non zero --> ");
39     }
40
41     fprintf(stderr, "managed to read: %zd characters\n", result);
42
43     printf("buffer contains: ");
44     int i;
45     for (int i = 0; i < result; ++i)
46         printf("_%c", buffer[i]);
47     printf("\n");
48
49     return 0;
50 }
```



- › You can send a signal to a running process using the kill system call function

```
#include <sys/types.h>
#include <signal.h>

int kill (pid_t pid, int sig);
```

Where pid is the process ID of the process to be signaled and sig is the signal to be sent.



Catching Signals

- › You can “catch” a signal by specifying a function that is called when the signal is received
- › This is done using the signal function:

```
typedef void (*sighandler_t)(int);  
  
#include <signal.h>  
                                         typedef of function prototype  
sighandler_t signal(int signum, sighandler_t handler);  
void (*signal(int sig, void (*catch)(int)))(int);
```

This complicated looking declaration means that signal is called with 2 arguments: the first is the signal to catch, the second is a pointer to the function that will be called when the signal is received. The signal function returns a pointer to the function that previously caught the signalphew.



Signal: Catch SIGINT

```
→ volatile int interrupted = 0;  
void impatient(int sigval) {  
    interrupted = 1;  
}
```

the signal handler (aka, impatient) should be as small as possible

volatile tells the compiler that the value of variable may change at anytime - without any action being taken by the code the compiler finds nearby

```
→ int main() {  
    signal(SIGINT, impatient);  
    printf("Now we wait...\n");  
    while (!interrupted)  
        usleep(10);  
    printf("Oh...you didn't like waiting\n");  
    printf("Program terminated\n");  
    return 0;  
}
```

Does it work?



THE UNIVERSITY OF
SYDNEY

errno





- › Most C functions report errors via return values, or their parameters
- › However, there is still an error reporting mechanism using a global variable called `errno`
- › Failed system calls typically set `errno` to be an integer value representing the type of error.
- › A companion function, `strerror` and `perror`, will print a textual description of the `errno` code.



- › The <errno.h> header file defines the integer variable

```
#include <errno.h>
#include <stdio.h>

int main() {
    FILE *fp = fopen("doesn't exist", "r");
    printf("errno: %d\n", errno);
    return 0;
}
```

- › **errno** is set by the **last** function call that will set **errno**.
- › There is only one **errno** value
- › It can be overridden by subsequent function calls
- › It is important to save this value immediately following

```
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>

void foo() {
    FILE *fp = fopen("doesn't exist", "r");
}

int main() {
    perror("first line of code");

    foo();
    printf("errno: %d\n", errno); // 2
    perror("we tried to read a file");

    errno = 0;

    int x = 5;
    printf("errno: %d\n", errno); // 0

    void *data = malloc(-1);
    printf("errno: %d\n", errno); // 12

    free(data);
    printf("errno: %d\n", errno); // 12

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>

void foo() {
    FILE *fp = fopen("doesn't exist", "r");
}

int main() {
    foo();
    printf("errno: %d\n", errno); // 2
    perror(""); // as you dont set errno to 0, even if the function works successfully, it still showing previous error msg

    int x = 5;
    printf("errno: %d\n", errno); // 2
    perror(""); // as you dont set errno to 0, even if the function works successfully, it still showing previous error msg

    void* data = (void*)malloc(1);
    printf("errno: %d\n", errno); // 2
    perror(""); // as you dont set errno to 0, even if the function works successfully, it still showing previous error msg

    free(data);
    printf("errno: %d\n", errno); // 2
    perror(""); // as you dont set errno to 0, even if the function works successfully, it still showing previous error msg

    FILE *fp2 = fopen("hi.txt", "r");
    printf("errno: %d\n", errno); // 2
    perror(""); // as you dont set errno to 0, even if the function works successfully, it still showing previous error msg

    return 0;
}
```

Low level file I/O

COMP2017/COMP9017





Low Level File Descriptors

- › Low level I/O is performed on file descriptors that are small integers indicating an open file
- › When a process is started file descriptor 0 is standard input, 1 is standard output, 2 is standard error output (UNIX)
- › System call functions operate on file descriptors



Low Level File Descriptors

- › low level I/O functions in C wrap system calls:
 - creat, open, close
 - read, write
 - ioctl io control
 - umask acess private data
- › eg read 100 characters from standard input into array “buffer”

```
ssize_t result = read(0, buffer, 100);
```

result is how many byte that has sucessful reading
negative value of ssize_t is error

```
$ man 2 open
```

“If I was to change anything in Unix it would be to spell creat with an e”

Ken Thompson



Low Level File Descriptors

- › **read()**

On error, -1 is returned, and `errno` is set appropriately. In this case, it is left unspecified whether the file position (if any) changes.

- › This may be interrupted by a signal. The way to check is to use `errno`

```
ssize_t result = read(...);  
if (result < 0) we save the value of errno to diagonal what has happened  
    error_val = errno;  
if (EINTR == error_val) // reattempt
```

- › These operations are blocking. There may be a need to interrupt them upon a new event.



Working with read and write

› Error checking

- errno is set to an error value
- signal can be sent by operating system

```
#include <errno.h>
...
signal(SIGINT, interrupted);
char buffer[100];
ssize_t result = read(0, buffer, 100);
// check for errors
int error_val = errno;
if (0 != error_val) {
    printf("read() was interrupted by signal\n");
}
```

Does it work?



Catching Signals

- › You can “catch” a signal by specifying a function that is called when the signal is received
- › This is done using the **sigaction()** function:

```
#include <signal.h>

int sigaction(int signum, const struct sigaction *act,
              struct sigaction *oldact); pass in memory address

struct sigaction {

    8 void      (*sa_handler) (int); function pointers
    8 void      (*sa_sigaction) (int, siginfo_t *, void *);
    4 sigset_t   sa_mask;
    4 int       sa_flags;
    8 void      (*sa_restorer) (void);

};
```

signal work differently on varies platforms, but
sigaction() works the same



Working with read and write

› Error checking

- errno is set to an error value
- signal can be sent by operating system

```
#include <errno.h>
...
// setup new handler
new_sig_int.sa_handler = interrupted;
new_sig_int.sa_flags = 0;
// install the new handler
sigaction(SIGINT, &new_sig_int, NULL); dont care about oldact

char buffer[100];
ssize_t result = read(0, buffer, 100);
// check for errors
int error_val = errno;
if (error_val != 0) {
    printf("read() was interrupted by signal\n");
}
It works
```



Low Level File Descriptors

- › Extra attention is needed when working with files at this level
 - Buffering
 - Sharing vs exclusive access (resource locking)
 - Errors and interruptions
 - Notifications (Linux)
 - Resource limit setting
 - Performance
- › `fcntl` - manipulate file descriptors
- › Valuable to have very fine control of file operations